

Intelligent Resume-Based Job Suggestion System

AI-powered job matching using AWS Bedrock,
MongoDB & Streamlit

By: Kanish Midhun

Problem Statement

- - Job seekers manually search and filter through hundreds of postings.
- - Recruiters struggle to match candidates at scale.
- - Resumes are unstructured and vary by format.
- - There is a need for an AI-powered, resume-aware job recommendation system.

Motivation & Use Cases

- Motivation:
 - Use LLMs and embeddings to understand resumes.
 - Reduce time to discover relevant jobs.
 - Provide transparent, explainable matching.
- Use Cases:
 - Fresh graduates looking for suitable roles.
 - Experienced professionals exploring targeted opportunities.
 - Career guidance platforms or university placement cells.

High-Level Architecture

Streamlit UI
(Resume Upload)

AWS S3
(Resume Storage)

Resume Lambda
Claude Parsing + Titan
Embeddings

Job Fetch Lambda
(JSearch API)

MongoDB Atlas
(resumes, jobs,
matches)

Job Matcher Lambda
Hybrid Ranking Engine

Streamlit UI
(Top Matches +
Heatmap)

Detailed Data Flow

- End-to-End Flow:
 - 1. User uploads resume via Streamlit.
 - 2. Resume stored in S3 → S3 event triggers Lambda.
 - 3. Lambda extracts text, calls Claude to parse structured JSON.
 - 4. Titan generates embeddings from parsed resume.
 - 5. Resume JSON + embeddings stored in MongoDB.
 - 6. Job Fetch Lambda queries JSearch API using inferred job title.
 - 7. Jobs enriched, deduplicated, embedded & stored in MongoDB.
 - 8. Job Matcher Lambda computes scores and writes final matches.
 - 9. Streamlit UI reads from MongoDB and displays ranked job list.

Resume Parsing Pipeline

- Steps:
- 1. S3 triggers Resume Lambda when a PDF is uploaded.
- 2. PyPDF2 extracts raw text from the resume.
- 3. Claude (via Bedrock) receives a structured parsing prompt.
- 4. Output is cleaned JSON with name, contact, skills, education, experience.
- 5. Titan embeddings are generated from the merged resume text.
- 6. Resume document + embedding stored in MongoDB (resumes collection).

Example Parsed Resume JSON (Conceptual)

- {
- "name": "Jane Doe",
- "email": "jane.doe@example.com",
- "skills": ["Python", "Data Analysis", "SQL"],
- "experience": [
 - {
 - "title": "Data Analyst",
 - "company": "ABC Corp",
 - "start_date": "2021-01",
 - "end_date": "2023-03"
 - }-],
- "education": [
 - {"degree": "B.Sc.", "field": "Computer Science"}-]
- }
- This JSON is produced by Claude from the raw resume text.

Job Fetching Pipeline (JSearch API)

- 1. Job Fetch Lambda reads the user's inferred job title from their parsed resume.
- 2. It calls the JSearch API (RapidAPI) with a query like 'nurse jobs in India'.
- 3. The API returns a list of job postings with title, company, description, and links.
- 4. Each job is enriched with a short summary and key skills via Claude.
- 5. Titan embeddings are computed for each job description.
- 6. Jobs are deduplicated and stored in MongoDB (jobs collection).

Example Job Record (Conceptual)

- {
- "title": "Registered Nurse - ICU",
- "company": "City Hospital",
- "location": "Chennai, India",
- "description": "Provide ICU patient care, monitor vitals, coordinate with doctors...",
- "job_link": "https://example.com/job123",
- "skills": ["ICU nursing", "Patient monitoring", "BLS", "ACLS"],
- "embedding": [... 1536-d vector ...],
- "created_at": "2025-11-01T10:00:00Z"
- }
- This enriched job data is used by the ranking engine.

Job Deduplication Logic

- To avoid repeated job postings, the pipeline uses:
 - - A hash of the job description text.
 - - Combination of (title + company + user_id).
 - - Job link uniqueness check.
- If any of these match an existing job record for the same user, the new job is skipped.

Hybrid Ranking Engine

- Final score is computed using:
- $\text{final_score} = 0.55 * \text{semantic_similarity}$
 - $+ 0.25 * \text{keyword_overlap}$
 - $+ 0.10 * \text{recency_weight}$
 - $+ 0.10 * \text{popularity_score}$
- semantic_similarity: cosine similarity between resume and job embeddings
- keyword_overlap: overlap between resume tokens and job tokens
- recency_weight: higher for recently fetched jobs
- popularity_score: based on metadata (e.g., apply link, employer rating)

Skill-Gap Detection & Explanation

- Claude is used again to compare:
 - Extracted resume skills
 - Required job skills
- It outputs:
 - missing_skills: list of skills the candidate does not mention
 - match_reason: short explanation why the job fits the candidate
- These are stored with each match and visualized as a skill-gap heatmap in Streamlit.

Streamlit User Interface

- The UI provides:
- - Resume upload (PDF) → upload to S3 + trigger pipeline.
- - 'Refresh Jobs' button → calls API Gateway → Job Fetch Lambda.
- - 'Load Matches' → reads MongoDB matches collection.
- - Top ranked jobs with scores and explanations.
- - Clickable job links to apply.
- - Skill-gap heatmap (matplotlib) .

Daily Auto Refresh (Optional)

- Design:
- 1. User enables 'Daily Refresh' from the UI (or is registered in a list).
- 2. A scheduled CloudWatch Event or external cron triggers Job Fetch Lambda daily.
- 3. For each registered user_id, it re-fetches and re-ranks jobs.
- 4. Updated matches are written to MongoDB.
- 5. The next time the user opens the UI, they see fresh recommendations.

Performance & Scalability

- Performance:
 - End-to-end flow can complete in a few seconds (depending on Bedrock & API latency).
 - Embeddings and ranking are efficient once data is in MongoDB.
- Scalability:
 - Lambda functions scale horizontally with parallel requests.
 - MongoDB Atlas can be scaled vertically and horizontally.
 - S3 is highly scalable for file storage.
 - Entire architecture is serverless and pay-per-use.

Error Handling & Robustness

- Key strategies:
- - Try/Except around external API calls (JSearch, Bedrock).
- - Logging to CloudWatch for each Lambda.
- - Fallback behaviors if Claude output is not valid JSON.
- - Timeouts and retries for network operations.
- - Graceful UI messages when no resume or matches are found.

Future Enhancements

- - Support multi-language resumes & job descriptions.
- - Integrate additional job boards and APIs.
- - Add user-specific preference learning (e.g., location, salary).
- - Build an admin dashboard for analytics.
- - Experiment with fine-tuned embedding models for specific domains.

Conclusion & Repository

- This project demonstrates:
 - - End-to-end AI job recommendation using LLMs and embeddings.
 - - A complete serverless pipeline on AWS.
 - - Integration of Bedrock, MongoDB, JSearch API, and Streamlit.
- GitHub Repository:
 - https://github.com/KanishMidhun/intelligent_resume_based_job_suggestion
- Thank you!